

Lingua Vertex: Language modeling with vertex ai

Abdul Vahed Shaik

San Jose State University, Department of Computer Software Engineering, San Jose, CA, USA

Abstract— The daily use of data for a variety of objectives, including marketing, business, sentiment analysis, election prediction, and communication has increased rapidly. In this paper I have implemented different language modeling tasks like text classification, sentiment analysis, and text summarization. I achieved this by fine-tuning the pretrained BERT model and leveraging GCP and Vertex AI to automate every stage of the ML system, including data collecting, model deployment, and serving.

Keywords— Text Classification, Summarization, MLOps, Sentiment Analysis, Question Answering, VertexAI, BERT, HuggingFace, Streamlit, Docker, TensorBoard, Pytorch.

I. INTRODUCTION

Text classification, Sentiment analysis, and summarization are crucial language modeling tasks. Language modeling is a tool for word prediction that analyses the structure of human language. The fundamental element of natural language processing (NLP), which can handle more text and language, is BERT. I have made use of this capability to carry out the language modeling duties by optimizing the pre-trained BERT model and automating all of the machine learning processes, such as collecting data, model deployment, and serving, through the use of Vertex AI and GCP.

I have performed the following language modeling related tasks using pre-trained models:

- 1.Text Classification
- 2.Text Summarization

II. RELATED WORK

Text is a widespread medium used by people to exchange ideas and have conversations. Many text papers are produced along with the endless interactions that take place on online sites. As a result, numerous academics have developed various Bert model variations in recent years that are simple to adjust for language modeling applications. Previous research on question answering was conducted by Talmar and Berant (2019)[11] who thoroughly examined question answering models using ten datasets, and by Yogatama et al. (2019)[10] who assessed a model across four datasets. DistilBert, which is quick, inexpensive, and trained by distilling Bert base, and the Stanford Question Answering dataset are used in this model.

Typical Deep Learning models for summarization use Seq2Seq architectures with RNNs (Nallapati et al.). Attention mechanisms frequently work in tandem with transformers (Vaswani et al., 2017; Lewis et al., 2020; Raffel et al., 2020a). To address the issue of these models not functioning effectively with financial information, HuggingFace developed the Pegasus summarizer model. With the help of this model and various user decoding techniques, users will be able to comprehend key terms in the text and produce a summary.DATASET

The datasets I utilized for the project, one for each of the language modeling tasks, are listed below:

1.Emotion Classification

Emotion from Hugging Face Datasets: Using this dataset, I finetuned the pretrained BERT model to better classify emotions. There are labels in the collection that belong to different classifications. Anger, fear, pleasure, love, sadness, and surprise are the six fundamental emotions represented in this collection of Twitter posts.

```
{
  "0": "Sadness",
  "1": "Joy",
  "2": "love",
  "3": "anger",
  "4": "fear"
}
```

Fig 1: Multiclass labels of the emotion classification

Below screenshot shows a preview of the dataset:

text (string)	label (class label)
i didnt feel humiliated	sadness
i can go from feeling so hopeless to so damned hopeful just from being around someone who cares...	sadness
im grabbing a minute to post i feel greedy wrong	anger
i am ever feeling nostalgic about the fireplace i will know that it is still on the property	love
i am feeling grouchy	anger
ive been feeling a little burdened lately wasnt sure why that was	sadness
ive been taking or milligrams or times recommended amount and ive fallen asleep a lot faster but i...	surprise
i feel as confused about life as a teenager or as jaded as a year old man	fear

Fig 2: Preview of the emotion dataset

III. METHODOLOGIES

A. Modeling

We The pre-trained BERT model has been refined and used to the emotion categorization job on Vertex AI. A large corpus of unlabeled text, comprising the whole Wikipedia (2,500 million words) and Book Corpus (800 million words), is used to pre-train the BERT model. Using BERT has the benefit of requiring little adjustment for the emotion classification task. Additionally, this leads to faster development and makes use of a smaller dataset overall.

B. MLOps

All of the necessary phases of an ML system, from data collection to model deployment and servicing, may be conducted and automated with the help of an ML pipeline.

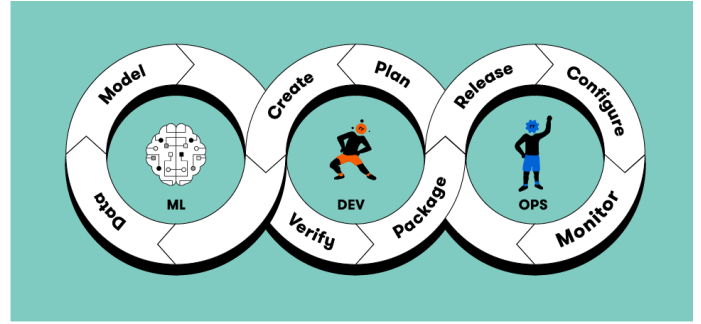


Fig 3: A typical MLOps architecture

The goal of MLOps, an ML engineering culture and practice, is to bring together ML system operation (Ops) and ML system development (Dev). By using MLOps, you provide automation and monitoring during the whole development of an ML system, from infrastructure management to integration, testing, releasing, and deployment. Given pertinent training data for their use case, we can put into practice and train an ML model with predictive performance on an offline holdout dataset. Nevertheless, creating an integrated machine learning system and keeping it running in production continually presents a greater difficulty than simply creating an ML model.

C. GCP and Vertex AI

I'm using the Vertex AI pipeline to custom train a pre-trained BERT model, adjust its hyperparameters, then use it to classify emotions.

Understanding Vertex AI:

Vertex AI is a unified MLOps platform designed to facilitate faster deployment, more experimentation, and confident model management for data scientists and ML engineers. With pre-trained and bespoke tooling within a single AI platform, it facilitates the faster development, deployment, and scaling of ML models.

It is an AI platform that helps in building a unified ML application by integrating Google cloud services. The key features of VertexAI includes:

- A unified UI for the entire ML workflow
- End-to-end integration for data and AI
- Pre-trained APIs for vision, video, natural language, and others
- Support for all opensource frameworks

Emotion Classification using Vertex AI:

Using the bert-base-cased model that was Pre-trained, I have developed an end-to-end training pipeline for the Vertex AI emotion categorization challenge. Emotion Classification is the dataset that is utilized in this task.

The primary goal of this work is to create, train, optimize, and implement a PyTorch model based on BERT on Vertex AI, with a focus on providing assistance for PyTorch model deployment and training on Vertex AI.

We start by creating notebook instance on the GCP console by running the following command on the cloud console:

```
gcloud notebooks instances create cmpe297-project --vm-image-project=deeplearning-platform-release --vm-image-family=pytorch-1-9-cu110-notebooks --machine-type=n1-standard-4 --location=us-central1-a --boot-disk-size=100 --accelerator-core-count=1 --accelerator-type=NVIDIA_TESLA_V100 --install-gpu-driver --network=default
```

Fig 4: Command to create a workbench instance on GCP

Upon running this command, a notebook instance gets created in the workbench section of GCP's vertex AI. The result can be visualized as follows:

[illegible]

Fig 5: Result of gcloud command execution for notebook creation

In order to create this notebook instance, we have upgraded the required quotas from 0 to 1 since running our code requires GPU.

Google Cloud notebooks cover all the necessary requirements to fulfill the training, fine tuning and deploying of our model for emotion classification. The requirements include:

- The Google Cloud SDK
- Git
- Python 3
- Virtualenv which is a tool to create an isolated Python environment.
- Jupyter notebook with Python 3
- Transformers
- Datasets
- Hypertune
- I have used ‘aiplatform’ library that simplifies ML workflows by utilizing wrapper classes and opinionated defaults.

The following APIs are required to be enabled as a requirement to run the project on GCP:



Fig 8: Enabling Vertex AI API



Fig 6: Enabling Cloud Storage API

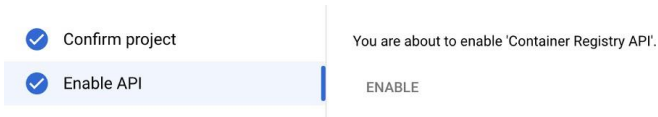


Fig 7: Enabling Container Registry API



Fig 8: Enabling Cloud Build API

Setting up the cloud storage bucket:

The model artifacts and the outcomes of the hyperparameter tuning are stored in the cloud storage bucket. This is used in the process of developing and implementing the model on Vertex AI. The following commands were run on the notebook instance in order to build a cloud bucket.

```
print(f"PROJECT_ID = {PROJECT_ID}")
print(f"BUCKET_NAME = {BUCKET_NAME}")
print(f"REGION = {REGION}")

! gsutil mb -l $REGION $BUCKET_NAME
```

Fig 9: Created a cloud bucket

Model Training

In this project, we use a pre-trained BERT model from Hugging Face Transformers to train the PyTorch emotion classification model locally first, and then we use the Vertex AI training service.

Training locally:

Using the Hugging Face Datasets library, I first downloaded the emotion dataset, which we then loaded using the 'load_dataset' function.

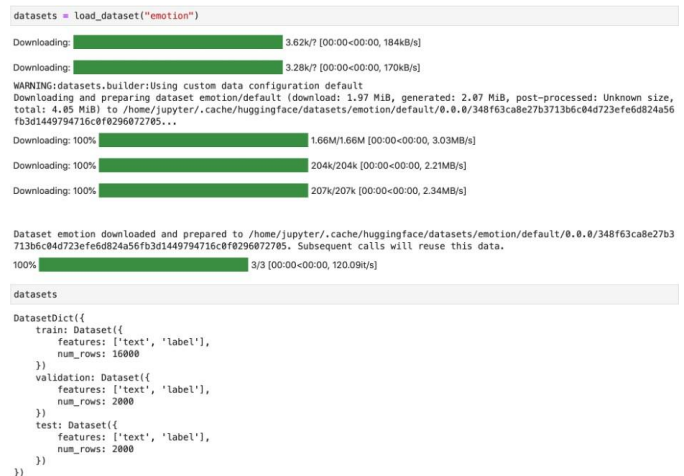


Fig 10: Loading the dataset on the notebook instance

In order to tokenize and preprocess the input data into the format needed by the model, I implemented the Hugging Face Transformers "Tokenizer" class. A screenshot of some of the tokens created for text input is shown below:

```
tokenizer("Hello and good morning!")

{'input_ids': [101, 8667, 1105, 1363, 2106, 106, 182], 'token_type_ids': [0, 0, 0, 0, 0, 0], 'attention_mask': [1, 1, 1, 1, 1, 1]}
```

Fig 11: Results of the tokenizer on input text

Model Fine-tuning

Transfer learning is a method that allows a deep learning model that was trained on a big dataset to function similarly on a smaller dataset. In order to accomplish this, the middle levels are frozen, and the input and output layers are changed to better fit our needs.

In this work, we are finetuning the pre-trained BERT model for our emotion categorization problem.



Fig 12: Fine-tuning BERT

A new head for emotion classification has been utilized in place of the one that was used to pre-train the model on masked language modeling tasks. When we fine-tune the model, pre-trained weights will be assigned; otherwise, they won't have any at first.

Then constructed a trainer object with all the necessary training settings, such as the batch size, learning rate, number of epochs, and weight decay. The same is seen in the screenshots below:

```
args = TrainingArguments(
    evaluation_strategy="epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=batch_size,
    per_device_eval_batch_size=batch_size,
    num_train_epochs=10,
    weight_decay=0.01,
    output_dir="/tmp/cls",
)
```

Fig 13: Creating a trainer object to fine-tune BERT

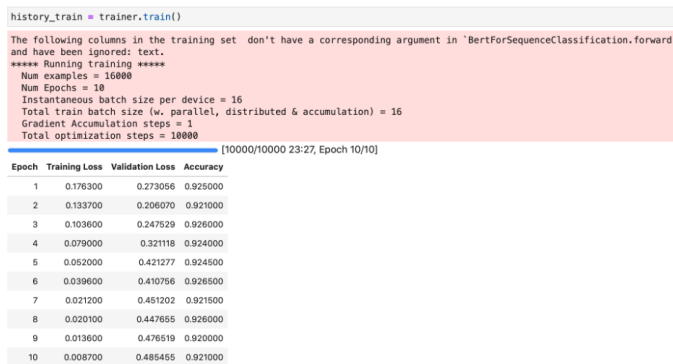


Fig 14: Results of fine-tuning BERT for 10 epochs

I also recorded TensorBoard logs for model training and noticed that fine tuning the BERT model for only one epoch suffices as it achieves good accuracy. I trained for 10 epochs so as to visualize the logs on TensorBoard as part of our study.

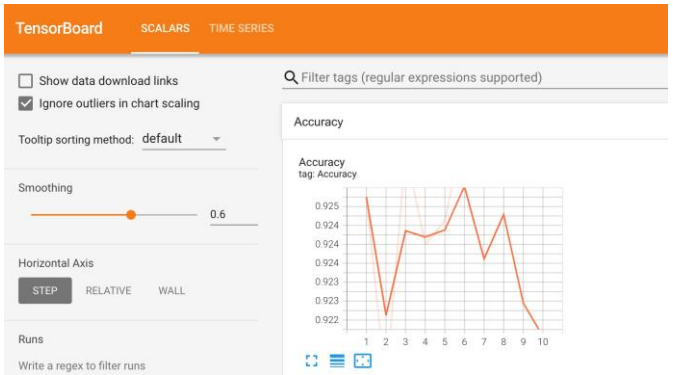


Fig 15: TensorBoard logs for tracking Accuracy



Fig 16: TensorBoard logs for tracking the training loss

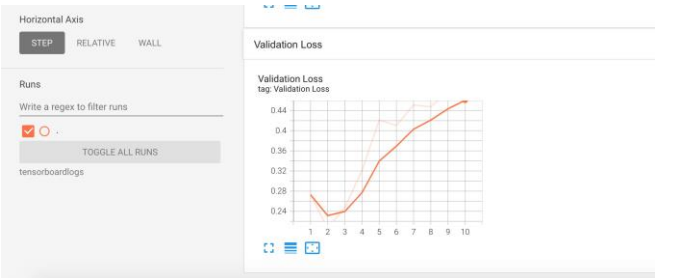


Fig 17: TensorBoard logs for tracking the validation loss

Additionally, we can observe that after training, the model artifacts are saved in the local folders:

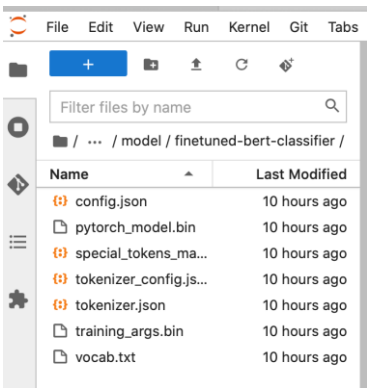


Fig 18: Model artifacts locally saved

Running predictions locally

After preprocessing the input, we can now predict the emotion labels for the text input.

```
# example #1
review_text = {
    "im feeling quite sad and sorry for myself but ill snap out of it soon"
}
predict_input = predict(review_text, saved_model_path)

loading configuration file https://huggingface.co/bert-base-cased/resolve/main/config.json from cache at /home/jupyter/.cache/huggingface/transformers/a883e8468a8fe98683bdc453f4fac622804f49de86d7cecae92365d4a0f829_a64a22196690ee8e2ead56f388a3ef3a58de9335926ccfa28618217db589387
Model config BertConfig {
  "architectures": [
    "BertForMaskedLM"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "position_embedding_type": "absolute",
  "transformers_version": "4.12.5",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 28996
}
```

Fig 19: Local predictions the fine-tuned BERT - 1

```
loading weights file ./models/pytorch_model.bin
All model checkpoint weights were used when initializing BertForSequenceClassification.

All the weights of BertForSequenceClassification were initialized from the model checkpoint at ./models.
If your task is similar to the task the model of the checkpoint was trained on, you can already use BertForSequenceClassification for predictions without further training.
Review text: im feeling quite sad and sorry for myself but ill snap out of it soon
Sentiment : Sadness
```

Fig 20: Local predictions the fine-tuned BERT - 2

transformers, datasets and tqdm – which are added in the setup.py file). We achieve this by creating a Python source with pre-built containers on Vertex AI.

We now run the custom training job on Vertex AI using Vertex SDK to create and submit the training job. The below screenshot shows the submission of the job from the notebook instance and also the training job successful on GCP's Vertex AI platform:

We can see that the model artifacts are successfully saved in the bucket after training.

Training on Vertex AI

Using Vertex AI to design a training pipeline is the most efficient way to handle larger datasets and models, like the ones we have. Vertex AI's training task is completed by packaging the code and setting up a training pipeline to coordinate training tasks. Our three steps have been as follows:

- Preparing the training code for distribution as a Python source code
- Using Vertex AI to finally deploy the optimal model to an endpoint;
- Hyperparameter training job.

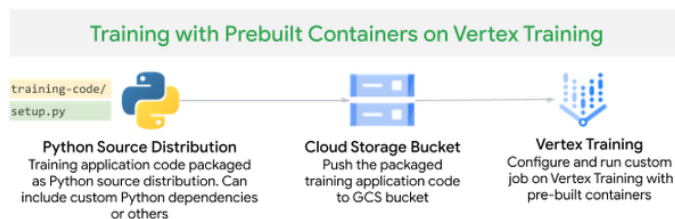


Fig 21: Custom training on Vertex AI

Step 1: Packaging the training application

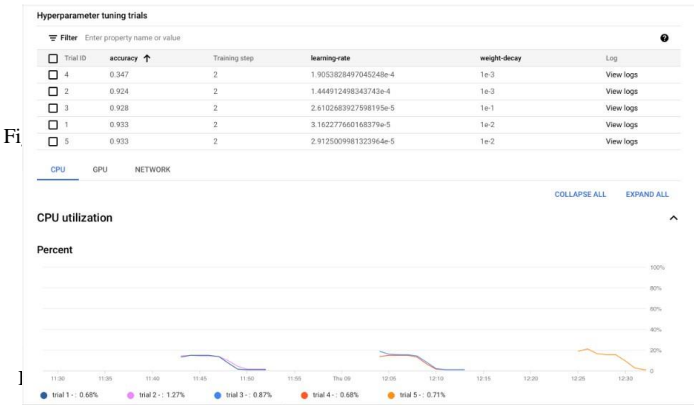
The first step involves packaging the training application code with all dependencies and uploading it to the Cloud Storage bucket or Container Registry (dependencies include

Hyperparameter Tuning

As I fine-tune the BERT model, I'm experimenting with hyperparameters like learning rate and weight decay. These hyperparameter values can significantly affect the model's performance because they are closely correlated with the training algorithm's behavior. I have packaged the training code and associated dependencies in a Docker container and used the Vertex Training service to automate the hyperparameter tuning. The container was subsequently pushed to Google Container Registry.

Following are the steps involved in performing a hyperparameter tuning job on Vertex AI service:

- 1. We repeatedly try the training application with the learning rate and weight decay hyperparameters set to their respective values.
- 2. Every trial run's outcomes are recorded by Vertex AI. Our training application uses the cloudml-hypertune Python tool to send the metrics to Vertex AI.
- 3. After the job is finished, we receive a Pandas data frame with a summary of all the trial runs that have the optimal settings based on all the criteria we requested.



We can visualize the logs for each of these training runs on Vertex AI. Below screenshot shows the training logs for Trial ID 2:

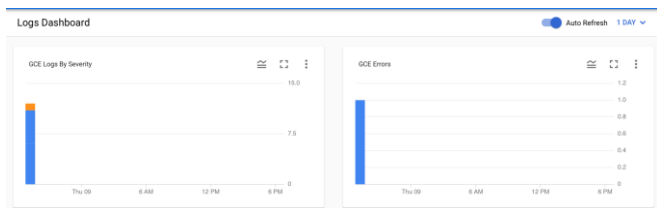


Fig 22: Hyperparameter tuning logs for one trial - 2

Deployment on VertexAI

A Docker container image running an HTTP server is used to deploy our optimized BERT model to VertexAI. The procedures for implementing the model for vertex predictions are as follows:

To service the model, I first constructed a customer container that was compatible with Vertex Predictions and bundled all the necessary model artifacts. At this point, I used the ping command to see how the container was doing. In order to confirm that the predictions are being served, I send a prediction request to the model on the container server upon success.

I then uploaded this model to Vertex as a model resource to serve predictions along with a customized container picture.

```
[174]: !curl http://localhost:7000/ping
{
  "status": "Healthy"
}

[177]: %bash -s $APP_NAME
APP_NAME=$1

cat > ./predictor/instances.json <<END
{
  "instances": [
    {
      "data": {
        "b64": "$echo 'I love pizza' | base64 --wrap=0"
      }
    }
  ]
}
END

curl -s -X POST \
-H "Content-Type: application/json; charset=utf-8" \
-d @./predictor/instances.json \
http://localhost:7000/predictions/$APP_NAME/
{"predictions": ["Joy"]}

[178]: !docker push $CUSTOM_PREDICTOR_IMAGE_URI
Using default tag: latest
The push refers to repository [gcr.io/starry-lens-333804/pytorch_predict_finetuned-bert-classifier]

7322eb6d: Preparing
f8c4de9e: Preparing
004756de: Preparing
855761fa: Preparing
abed59b2: Preparing
782ec13d: Preparing
5632bd30: Preparing
58eacbf1: Preparing
```

Fig 23: Pushing the best performing model after hyperparameter tuning

Additionally, I developed a Vertex endpoint and implemented our model resource to provide emotion predictions predicated on the text input.

IV. EXPERIMENTS

Text Summarization

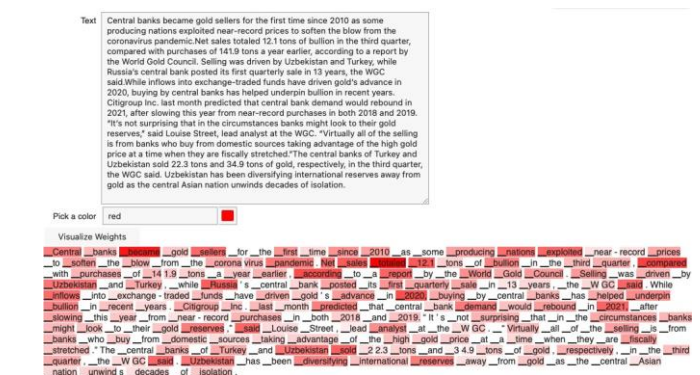
As a language modeling task, summarization compresses or summarizes a text document by retaining only the most important details, hence assisting in preventing information overload.

I have utilized Hugging Face's Pegasus model, which is specifically designed for financial statement summaries, to execute summarization. Bloomberg market and financial news was used to train it.

Given the assumption that a term with a high word density is more significant for summarization, I highlight the words based on their self-attention weights: high-weight words are bolded, while low-weight words are highlighted in a lighter manner.

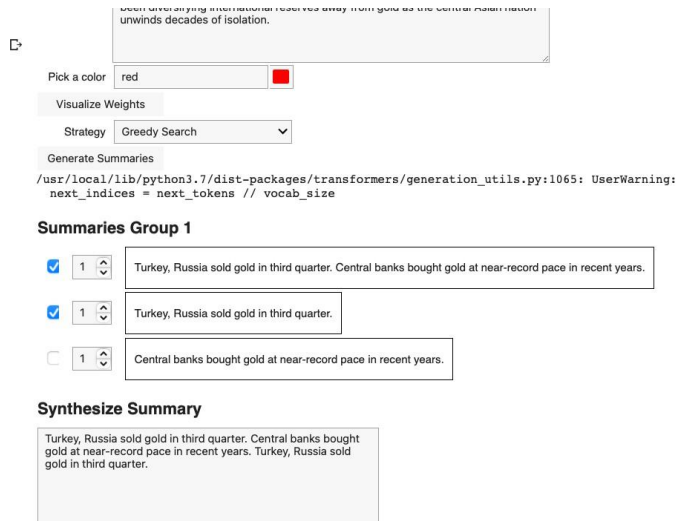
We The user can select from a variety of decoding algorithms to generate summaries, and based on the decoding strategy chosen, a control to adjust various hyperparameters is offered. We can select which of the several different

summaries produced by the model to combine into a comprehensive overview.



The different decoding strategies that can be selected are *Greedy Search*, *Beam Search*, *Diverse Beam Search*, and *Top k sampling*.

Greedy search selects the token with the highest probability at each time step.



Beam search chooses a number of tokens with the highest probability based on the beam width in addition to the token with the highest probability at each time step. The beam width and the number of final generated beams are equal.

Visualize Weights

StrategyBeam Search

#Beams5

Sequences2

Generate Summaries

/usr/local/lib/python3.7/dist-packages/transformers/generation_utils.py:1065: UserWarning: __floordiv__ is deprecated
next_indices = next_tokens // vocab_size

Summaries Group 1

☐

1

Turkey, Russia sold gold in third quarter. Central banks bought gold in recent years to cushion blow from pandemic.

☒

1

Turkey, Russia sold gold in third quarter.

☐

1

Central banks bought gold in recent years to cushion blow from pandemic.

Summaries Group 2

☐

1

Turkey, Russia sold gold in third quarter. Central banks bought bullion in recent years to cushion blow from pandemic.

☐

1

Turkey, Russia sold gold in third quarter.

☒

1

Central banks bought bullion in recent years to cushion blow from pandemic.

Synthesize Summary

Turkey, Russia sold gold in third quarter. Central banks bought bullion in recent years to cushion blow from pandemic.

Diverse beam search follows almost the beam search algorithm, but also adds a diversity penalty in order to enhance the diversity between the top most probable generated beams.

Uzbekistan sold 22.3 tons and 34.9 tons of gold, respectively, in the third quarter, the WGC said. Uzbekistan has been diversifying international reserves away from gold as the central Asian nation unwinds decades of isolation.

Pick a colorred

Visualize Weights

StrategyDiverse Beam Search

#Beams10

#Groups5

Penalty2.30

Sequences1

Generate Summaries

Summaries Group 1

☒

1

Turkey, Russia sold gold in third quarter. Central banks bought gold in recent years to cushion blow from pandemic.

☐

1

Turkey, Russia sold gold in third quarter.

☒

1

Central banks bought gold in recent years to cushion blow from pandemic.

Synthesize Summary

Turkey, Russia sold gold in third quarter. Central banks bought gold in recent years to cushion blow from pandemic. Central banks bought gold in recent years to cushion blow from pandemic.

With **top-k sampling**, the range of potential future tokens is restricted to the top k higher-ranked tokens in the distribution.

iscally stretched. The central banks of Turkey and Uzbekistan sold 22.3 tons and 34.9 tons of gold, respectively, in the third quarter, the WGC said. Uzbekistan has been diversifying international reserves away from gold as the central Asian nation unwinds decades of isolation.

Pick a colorred

Visualize Weights

StrategyTop-k Sampling

Top-k50

Sequences1

Generate Summaries

/usr/local/lib/python3.7/dist-packages/transformers/generation_utils.py:1242: UserWarning
next_indices = next_tokens // vocab_size

Summaries Group 1

☒

1

Turkey, Russia sold gold in third quarter. Central banks bought gold in recent years to cushion blow from pan

☒

1

Turkey, Russia sold gold in third quarter.

☐

1

Central banks bought gold in recent years to cushion blow from pandemic.

Synthesize Summary

Turkey, Russia sold gold in third quarter. Central banks bought

V. EVALUATION AND RESULTS

To display our model predictions on the user interface, I created a Streamlit application. By creating a service account key on the cloud console, authenticated the application to call the endpoint after deploying the model to the Vertex AI endpoint.

Subsequently used Google Cloud's "aiplatform" package to offer the fine tuned BERT model predictions to the Streamlit application..

Below are the screenshots of the model serving predictions for each of the class labels in the emotion dataset (sadness (0), joy (1), love (2), anger (3), fear (4)).

Advanced Deep Learning Emotion classification

Model is finetuned on pretrained bert

Type in your text here

I am pissed at myself for no reason

Run

anger



Advanced Deep Learning Emotion classification

Model is finetuned on pretrained bert

Type in your text here

I love pizzas

Run

Joy



Advanced Deep Learning Emotion classification

Model is finetuned on pretrained bert

Type in your text here

I am feeling bad about Puneeth's demise

Run

Sadness



Advanced Deep Learning Emotion classification

Model is finetuned on pretrained bert

Type in your text here

I am scared of exams

Run

fear



Advanced Deep Learning Emotion classification

Model is finetuned on pretrained bert

Type in your text here

a gentle tingle throughout almost as if i was feeling the healing taking place at a cellular level

Run

love



VI. CHALLENGES

One Model Training: Since it is a Large language model, it was a quite time consuming to train and analyze the model.

Hyperparameter tuning: Finding the right set of hyperparameters was a bit challenging since I was required to run multiple trials of them.

VII. CONCLUSION

Because of the large size of our model (pre-trained BERT) and dataset (Huggingface's Emotion dataset), Vertex AI's distributed server design and processing capability made establishing a training pipeline the most efficient approach. I was able to develop a repeatable pipeline and adjust the hyperparameters to determine the optimal model thanks to Vertex AI. With Vertex AI, we were able to effectively custom train our BERT model and deploy it to the endpoint. By utilizing the Streamlit application's model API, we were also able to receive the predictions effectively.

Additionally, we used pre-trained BERT models to perform language modeling tasks like Text Summarization on several datasets, and we achieved noteworthy model results. Pre-trained models enable us to create more quickly, use less data, and produce better outcomes.

Through this project, I've deepened my understanding of language models, their capabilities, and limitations. I've also learned about training data, bias mitigation, and model interpretability.

REFERENCES

- [1] <https://huggingface.co/datasets/emotion>
- [2] Li, Q., Peng, H., Li, J., Xia, C., Yang, R., Sun, L., ... & He, L. (2020). A survey on text classification: From shallow to deep learning. arXiv preprint arXiv:2008.00364.
- [3] <https://github.com/nlpunibo/Question-Answering-SQUAD>
- [4] <https://huggingface.co/transformers/>
- [5] <https://cloud.google.com/vertex-ai/docs/tutorials/image-recognition-auto>
- [6] Pegasus model:
<https://huggingface.co/human-centered-summarization/financial-summarization-pegasus>